Who wins? Quality and speed are different contests

2,848 graph/budget comparisons; equal weight per comparison. Quality ties split one point.

Native solver defaults, not tuned per instance. Mean gaps expose failures that median gaps can hide.

The budget map

Columns: solver. Each panel: runs x integration steps. All instances are included.

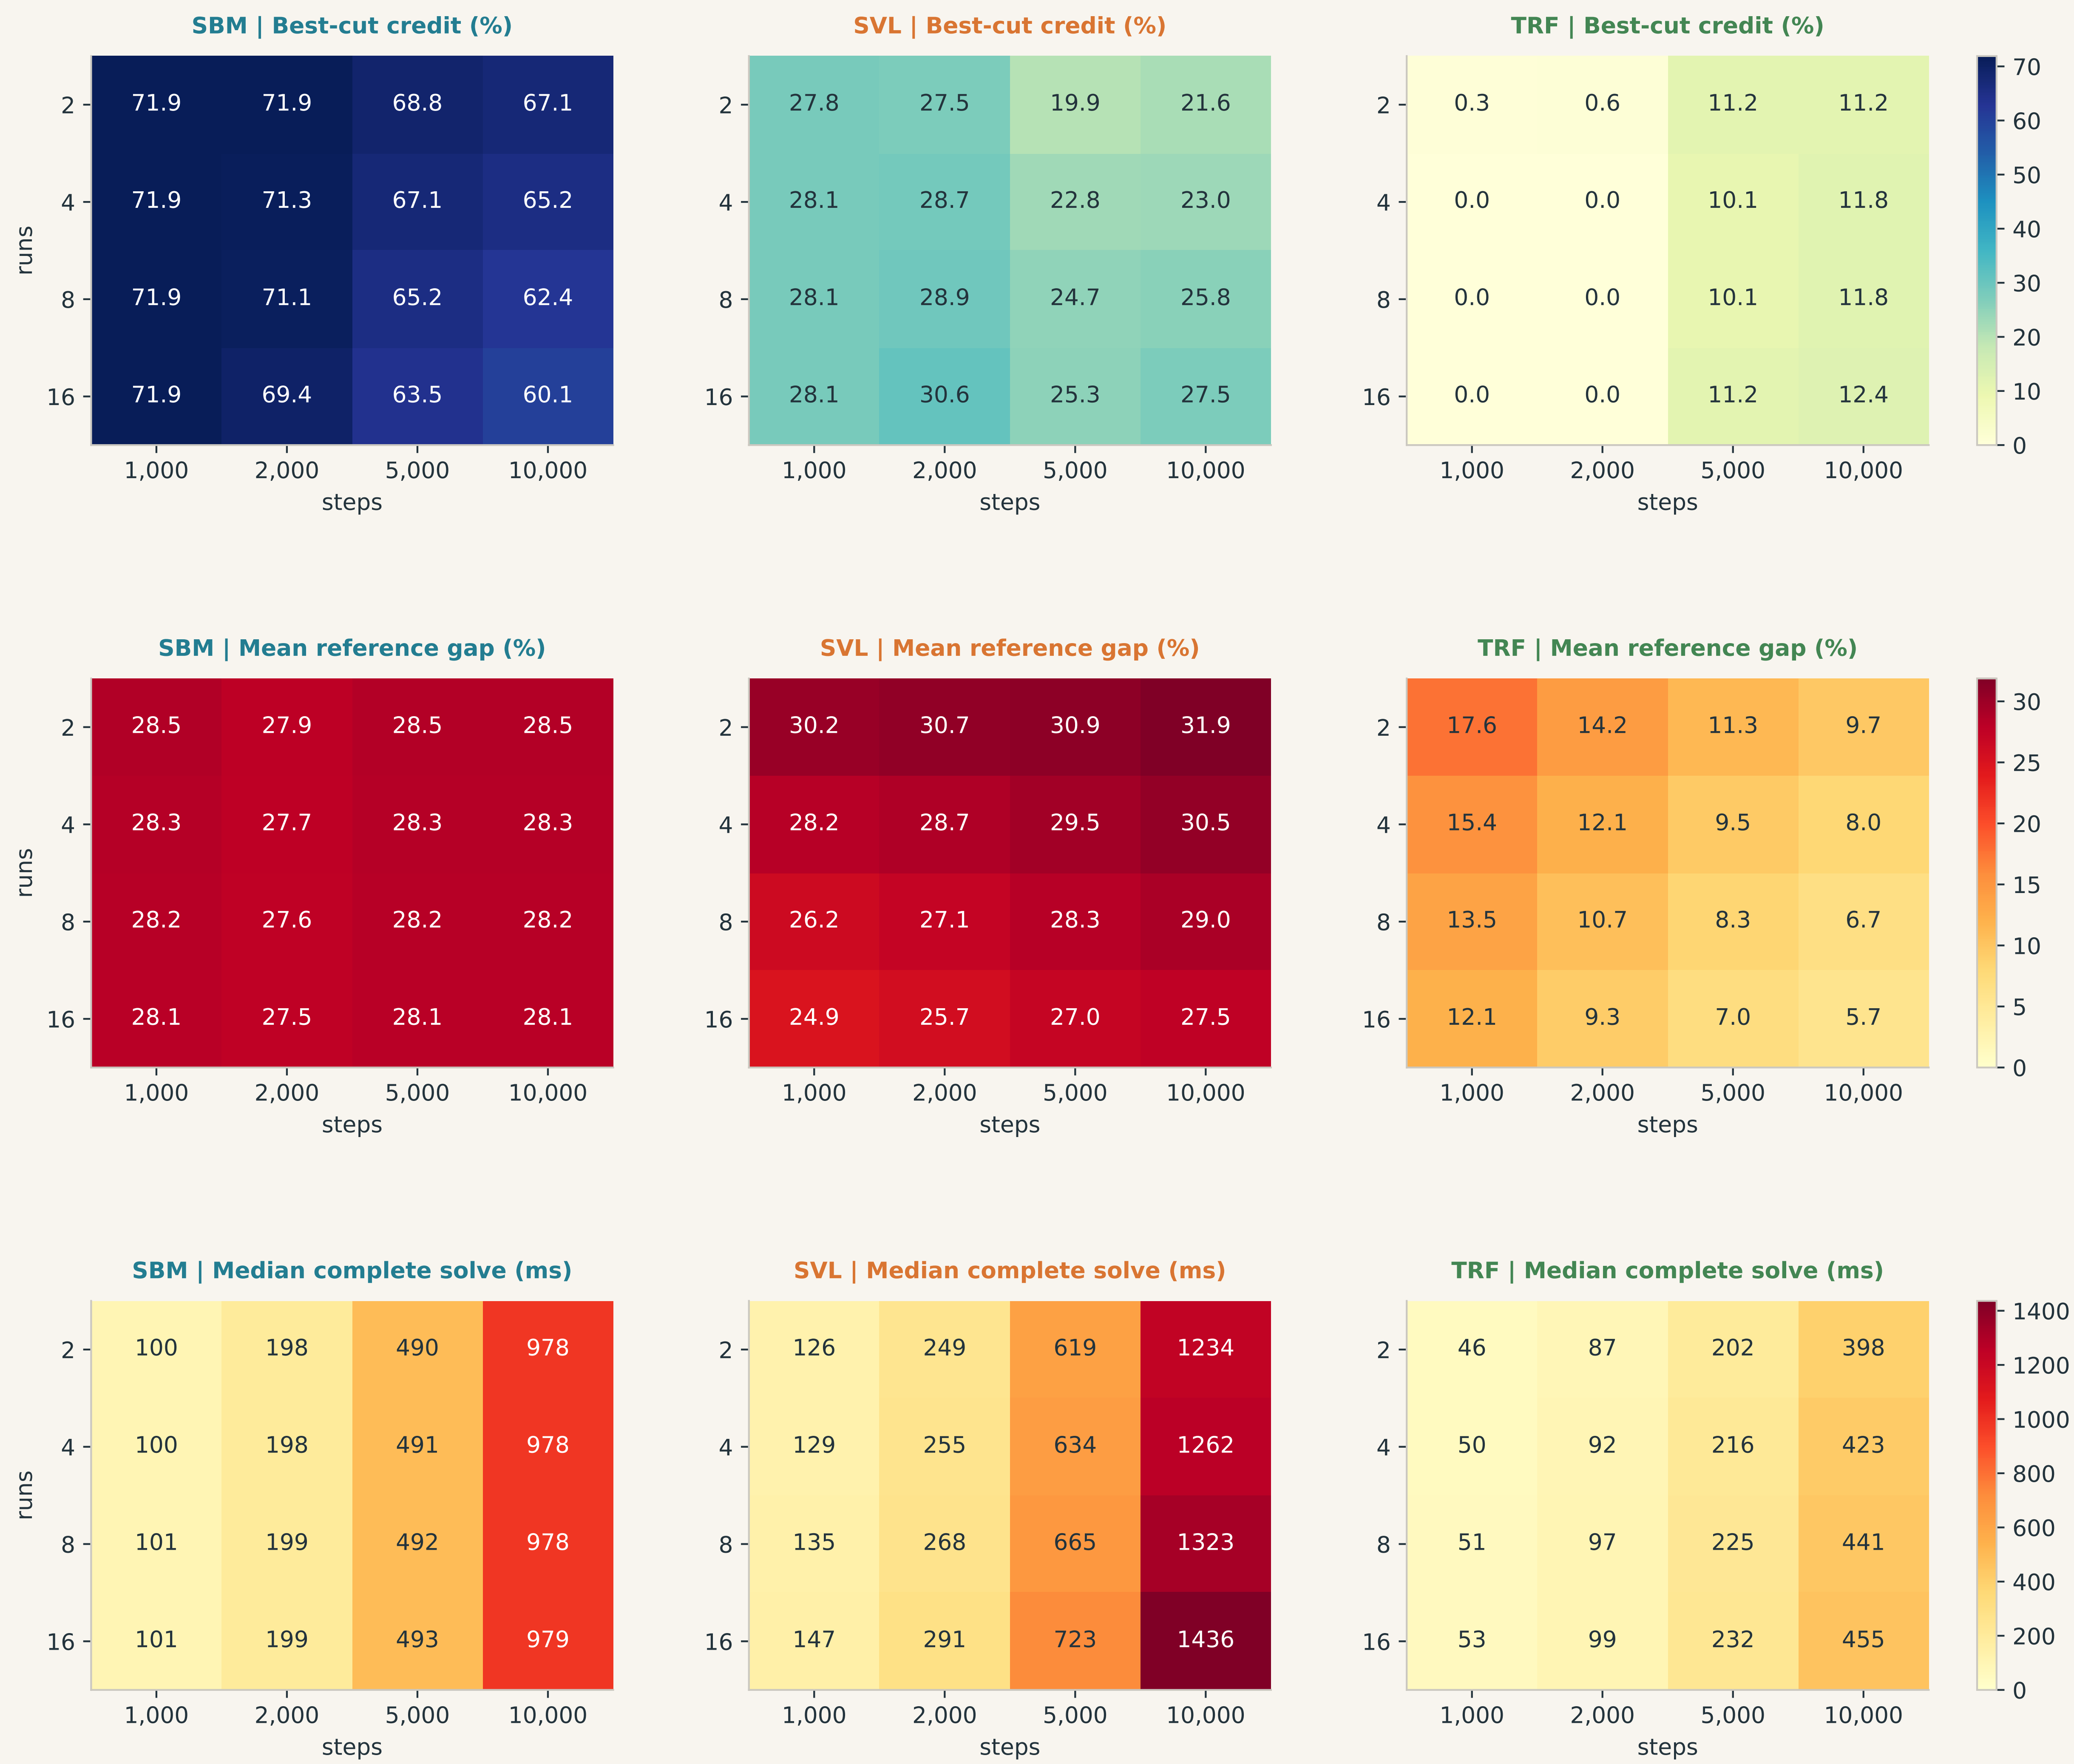

More blue = more quality wins. More red = larger error / longer runtime. Each row uses one shared color scale.

Speed versus solution quality

Each point aggregates all instances at one budget. Lower-left is better; labels show steps.

Solve latency is NOT time-to-first-target. No trajectory hit times or independent-seed success probabilities were recorded.

Look beyond the averages

Empirical distributions over all graph/budget comparisons. A curve farther left is better.

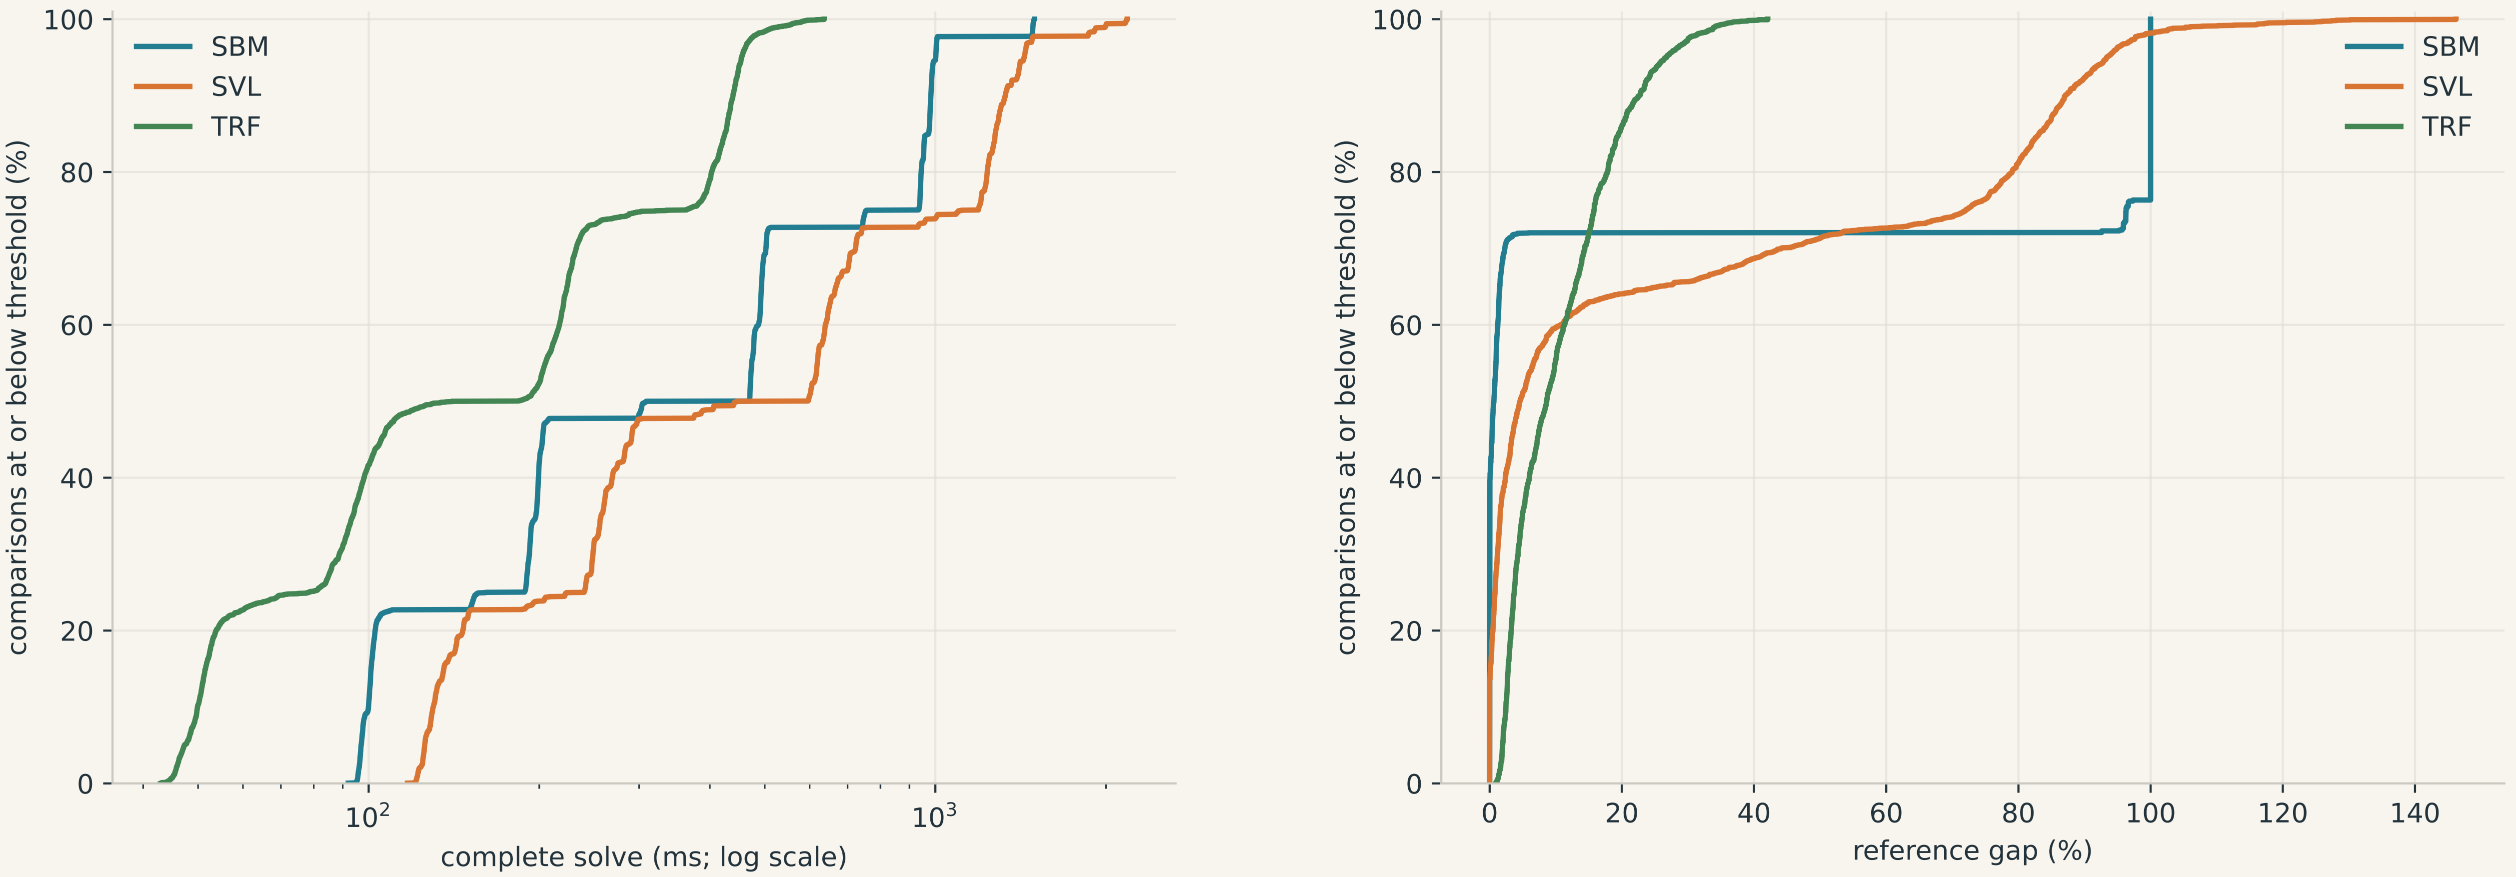

Timings: median of three same-seed trials per solve. These distributions mix graph families and sizes, not confidence intervals.

Exactly where each solver wins

One row per graph; one column per runs/steps setting. Zoom into the PNG or PDF for graph labels.