



Who wins? Quality and speed are different contests

2,055 graph/budget comparisons; equal weight per comparison. Quality ties split one point.


Share of best-cut credit




Fastest complete solve



Published reference matched












Mean shortfall from reference



Native solver defaults, not tuned per instance. Mean gaps expose failures that median gaps can hide.

The budget map






Available complete comparisons only; n = graphs per setting. Grey = no completed results. Cohorts can differ.



More blue = more quality wins. More red = larger error / longer runtime. Each row uses one shared color scale.

Speed versus solution quality



Each point aggregates available graphs at one budget; cohorts may differ. Lower-left is better; labels show steps.



Solve latency is NOT time-to-first-target. No trajectory hit times or independent-seed success probabilities were recorded.

Look beyond the averages

Empirical distributions over all graph/budget comparisons. A curve farther left is better.



Timings: median of three same-seed trials per solve. These distributions mix graph families and sizes, not confidence intervals.

Exactly where each solver wins

One row per graph; one column per runs/steps setting. Zoom into the PNG or PDF for graph labels.

